

# Report

## Methodology :



## Results :

	Model	Number of Conv Layers	Batch Normalization	Weight Decay	Learning Rate Scheduler	Learning Rate	optimizer	Loss Function	Accuracy	precision	recall	f1-score	Number of Epochs
0	First Model	3	No	No	No	0.001	Adam	CrossEntropyLoss	0.79	0.82	0.83	0.83	20
1	Second Model	4	Yes	Yes	Yes	0.001	Adam	CrossEntropyLoss	0.86	0.88	0.90	0.89	50
2	Third Model	4	Yes	Yes	Yes	0.001	Adam	CrossEntropyLoss	0.87	0.91	0.92	0.91	100

## Preprocessing :

- **normalized** the data using `transforms.Normalize` by 0.5
- **scaled** using `transforms.ToTensor()` scales the pixel values from the range [0, 255] to [0, 1].
- **Agumention** using **Random Affine Transformation,Random Horizontal Flip,Random Affine Transformation**

```
[17] transform_train = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomAffine(degrees=0, translate=(0.05, 0.05)),
    transforms.RandomAffine(degrees=0, scale=(0.95, 1.05)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])
dataset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform_train)
```

## Loading and splitting the data by 0.2 validation split with Applying The preprocessing steps

```
(validation_split = 0.2
dataset_size = len(dataset)
split = int(validation_split * dataset_size)
train_size = dataset_size - split)

[ ] train_data, valid_data = torch.utils.data.random_split(dataset, [train_size, split])
train_loader = torch.utils.data.DataLoader(train_data, batch_size=128, shuffle=True, pin_memory=True)
valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=128, shuffle=False, pin_memory=True)

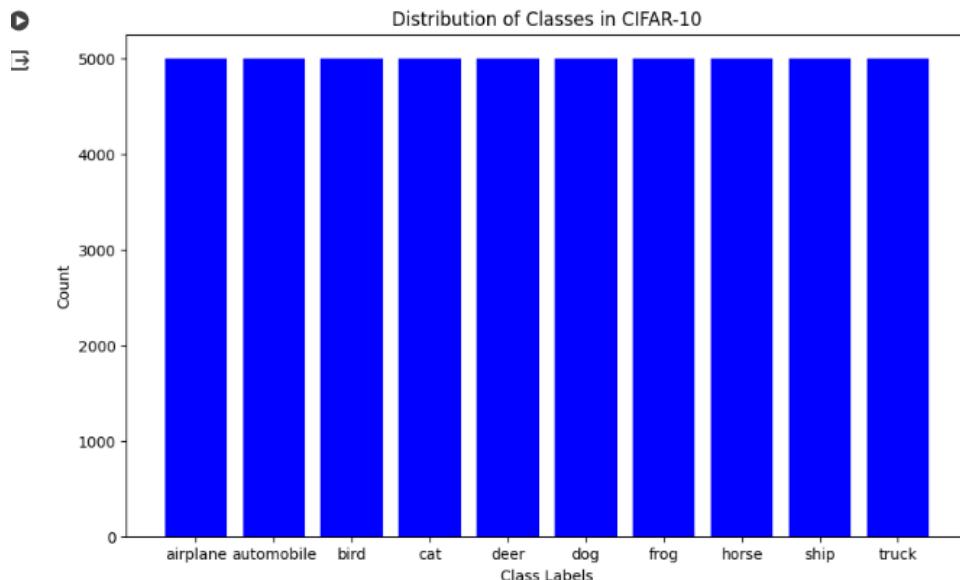
test_data = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform_test)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=128, shuffle=True, pin_memory=True)
```

```
▶ print(dataset)
└─ Dataset CIFAR10
    Number of datapoints: 50000
    Root location: ./data
    Split: Train
    StandardTransform
    Transform: Compose(
        RandomHorizontalFlip(p=0.5)
        RandomAffine(degrees=[0.0, 0.0], translate=(0.05, 0.05))
        RandomAffine(degrees=[0.0, 0.0], scale=(0.95, 1.05))
        ToTensor()
        Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
    )
```

```
▶ print(test_data)
Dataset CIFAR10
Number of datapoints: 10000
Root location: ./data
Split: Test
StandardTransform
Transform: Compose(
    ToTensor()
    Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
)
```

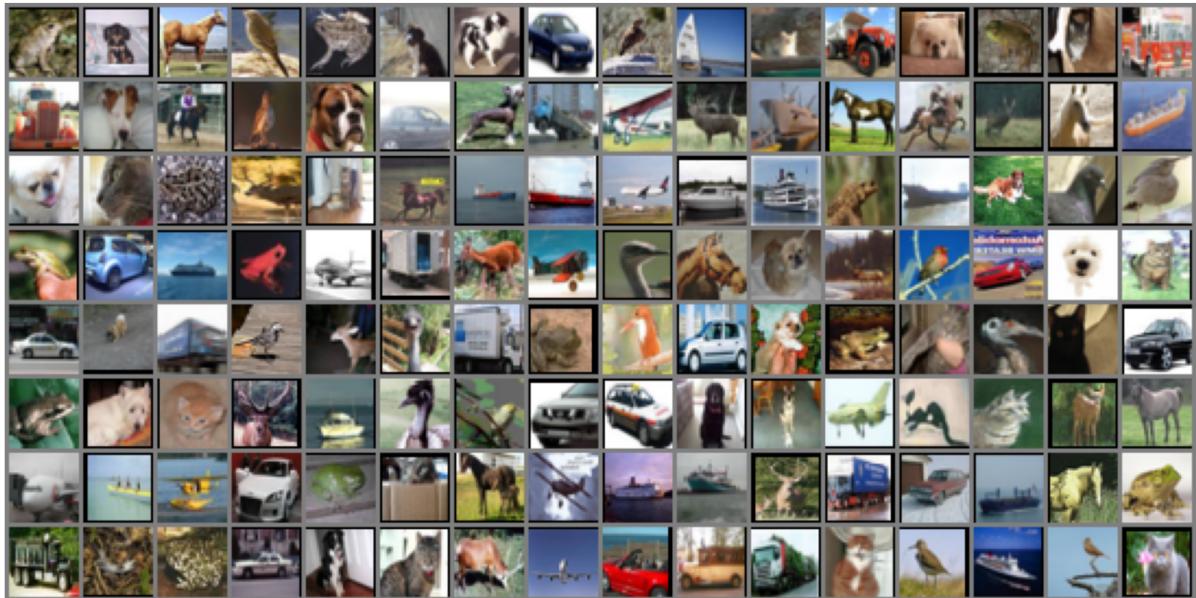
## Classes of the dataset

- we have 10 balanced classes

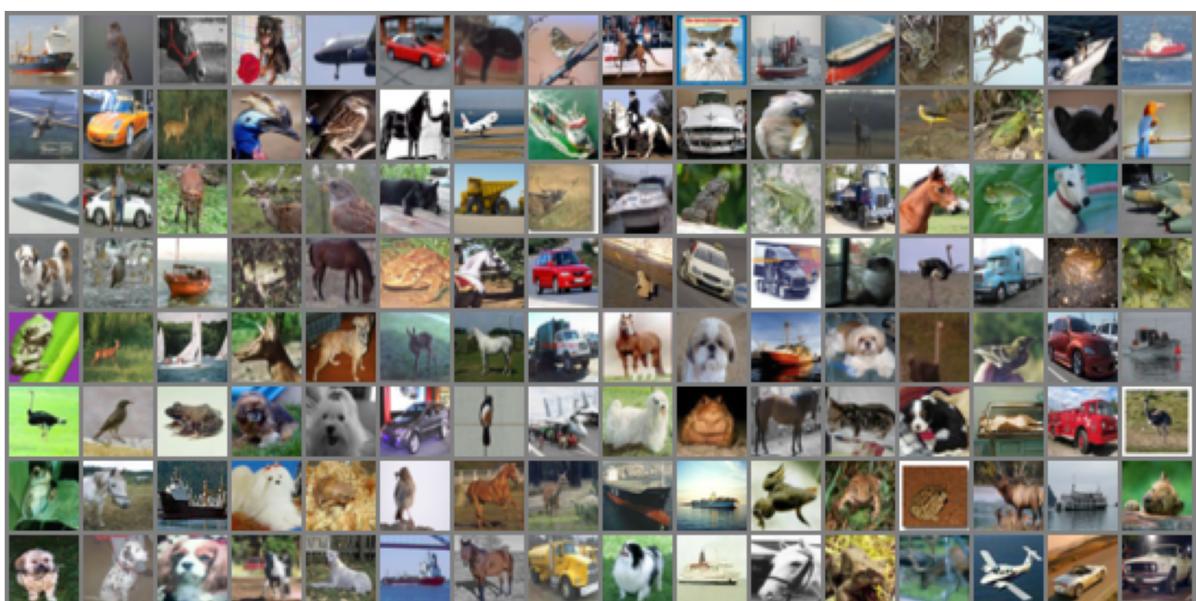


## Visualization:

## images of the train dataset

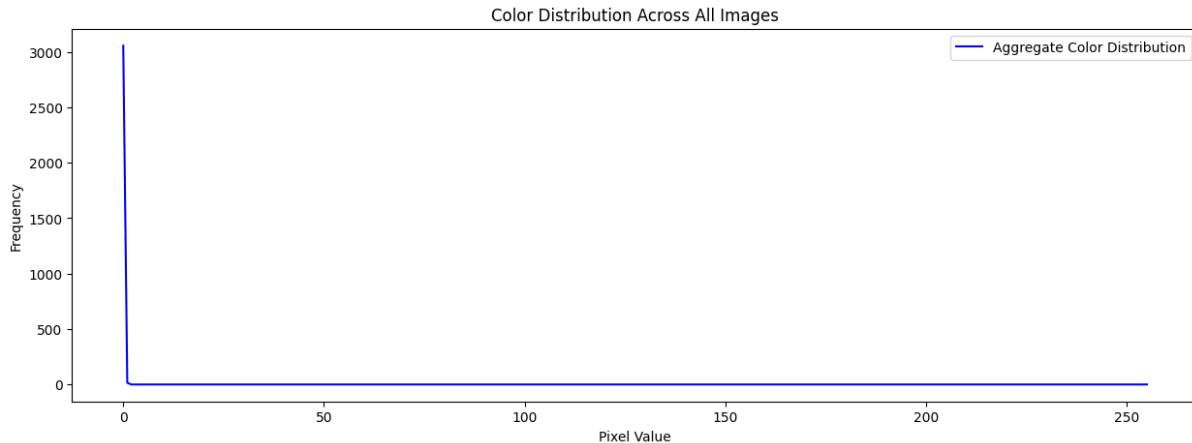


## **images of the test dataset :**



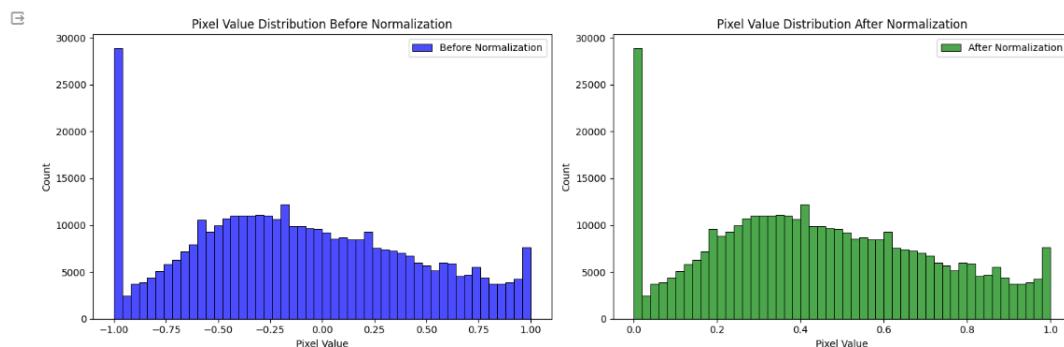
## Color distribution across all images in a dataset Graph :

It helps to understand the prevalent colors and intensities present in the images as I was trying to know more about the images in the dataset



## Pixel Value Distribution before and after normalization :

Comparing distributions show the impact of normalization on the data range and to be sure that the normalization is applied Correctly



## Training and Evaluating and testing Function

```
def train_and_evaluate(model, train_loader, valid_loader, test_loader):
    train_losses = []
    test_losses = []
    valid_accuracies = []
    all_labels = []
    all_predictions = []
```

```
for epoch in range(epochs):
    model.train()
    running_loss = 0.0

    for inputs, labels in train_loader:
        if train_on_gpu:
            inputs, labels = inputs.cuda(), labels.cuda()

            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()

    avg_train_loss = running_loss / len(train_loader)
    train_losses.append(avg_train_loss)

model.eval()
correct = 0
total = 0

with torch.no_grad():
    for inputs, labels in valid_loader:
        if train_on_gpu:
            inputs, labels = inputs.cuda(), labels.cuda()

            outputs = model(inputs)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
            all_labels.extend(labels.cpu().numpy())
            all_predictions.extend(predicted.cpu().numpy())

accuracy = correct / total
valid_accuracies.append(accuracy)
```

```
print(f'Epoch {epoch + 1}/{epochs}, Training Loss: {train_loss:.4f}\n')

model.eval()
test_loss = 0.0

with torch.no_grad():
    for inputs, labels in test_loader:
        if train_on_gpu:
            inputs, labels = inputs.cuda(), labels.cuda()

        outputs = model(inputs)
        loss = criterion(outputs, labels)
        test_loss += loss.item()

avg_test_loss = test_loss / len(test_loader)
test_losses.append(avg_test_loss)

print(f'Epoch {epoch + 1}/{epochs}, Training Loss: {avg_train_loss:.4f}, Test Loss: {avg_test_loss:.4f}\n')

return train_losses, test_losses, valid_accuracies, all_predictions
```

## First Model :

```
▶ class ImprovedNet(nn.Module):
    def __init__(self):
        super(ImprovedNet, self).__init__()
        self.conv1 = nn.Conv2d(3, 64, 3, padding=1)
        self.conv2 = nn.Conv2d(64, 128, 3, padding=1)
        self.conv3 = nn.Conv2d(128, 256, 3, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(256 * 4 * 4, 512)
        self.fc2 = nn.Linear(512, 256)
        self.fc3 = nn.Linear(256, 10)
        self.dropout = nn.Dropout(0.5)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        x = x.view(-1, 256 * 4 * 4)
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = F.relu(self.fc2(x))
        x = self.dropout(x)
        x = self.fc3(x)
        return x

improved_model = ImprovedNet()
if train_on_gpu:
    improved_model.cuda()

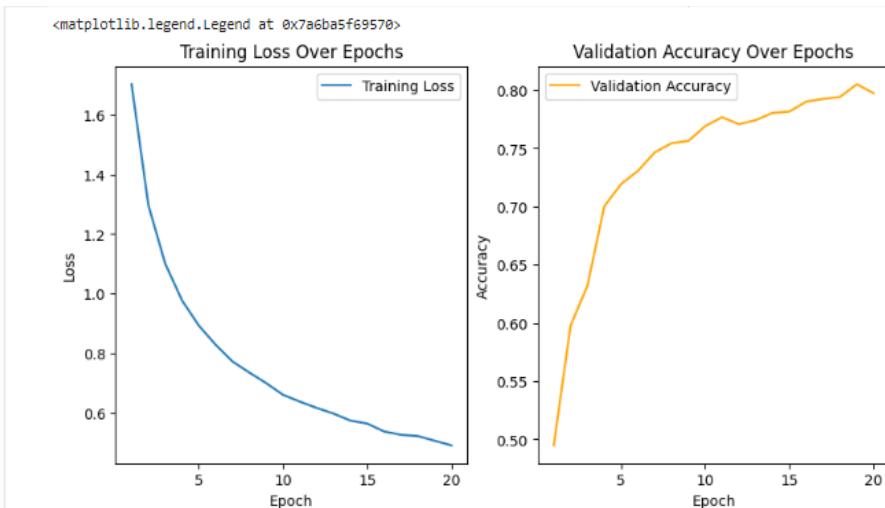
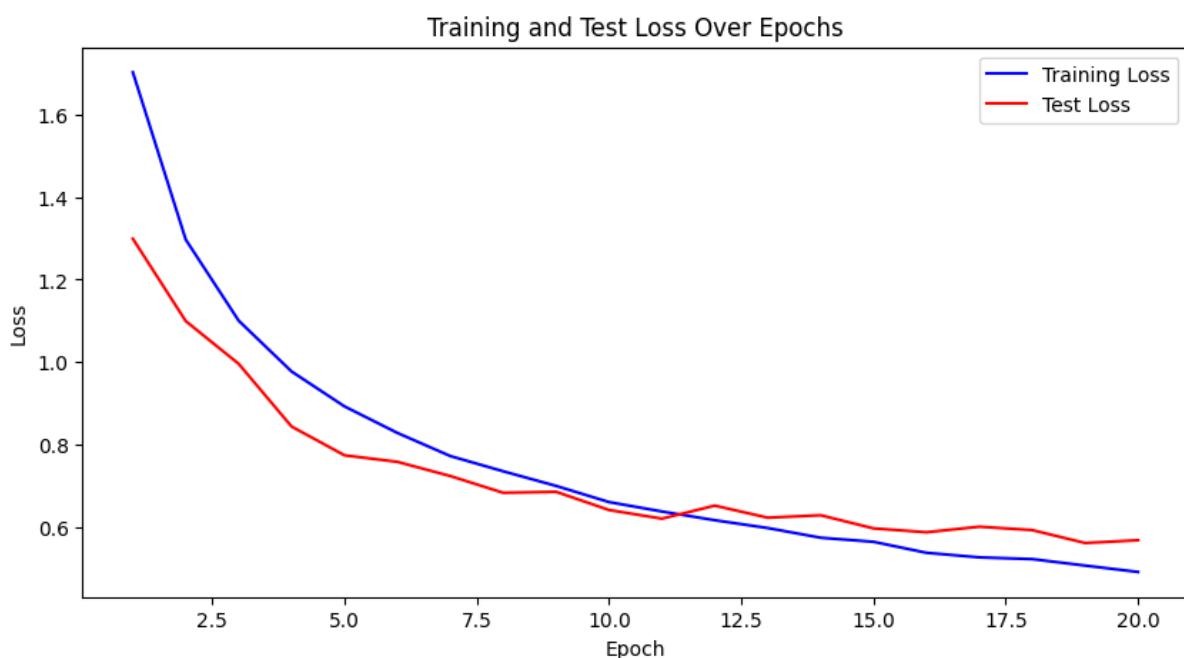
[ ] criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(improved_model.parameters(), lr=0.001)
```

## Confusion Matrix and Classification Report :

		Confusion Matrix									
		Confusion Matrix									
Actual	Predicted	0	1	2	3	4	5	6	7	8	9
		14562	327	1323	403	343	117	142	235	995	633
		332	16476	85	106	34	81	160	77	515	1652
		749	77	11437	1054	1102	861	699	443	208	83
		393	90	1129	10719	752	3404	957	730	264	262
		349	24	2248	1242	14242	987	973	1204	104	63
		123	50	1088	3505	518	13628	356	1196	77	68
		182	145	1339	1753	1114	553	16364	166	108	198
		185	66	576	602	1469	1109	78	15874	57	295
		1544	507	295	403	142	109	208	126	17011	426
		661	1358	120	333	104	171	123	209	521	17340

Classification Report:				
	precision	recall	f1-score	support
0	0.76	0.76	0.76	19080
1	0.86	0.84	0.85	19518
2	0.58	0.68	0.63	16713
3	0.53	0.57	0.55	18700
4	0.72	0.66	0.69	21436
5	0.65	0.66	0.65	20609
6	0.82	0.75	0.78	21922
7	0.78	0.78	0.78	20311
8	0.86	0.82	0.84	20771
9	0.82	0.83	0.83	20940
accuracy			0.74	200000
macro avg	0.74	0.74	0.74	200000
weighted avg	0.74	0.74	0.74	200000

## Evaluation Graphs



## The changes I made to get 2nd model :

- I added learning rate scheduler For convergence speed
- I added new Conv layer and Batch Normalization
- I added weight\_decay =1e-4 as a form of L2 regularization applied to the optimizer to help prevent overfitting.
- I added lr\_scheduler to speed convergence
- I changed the number of epochs to be 50 instead of 20 ( I actually tried 30 and I applied some of the changes above but i got (80,82,83) as Accuracy so I kept the best of them which was the below model

## Second Model Arch :

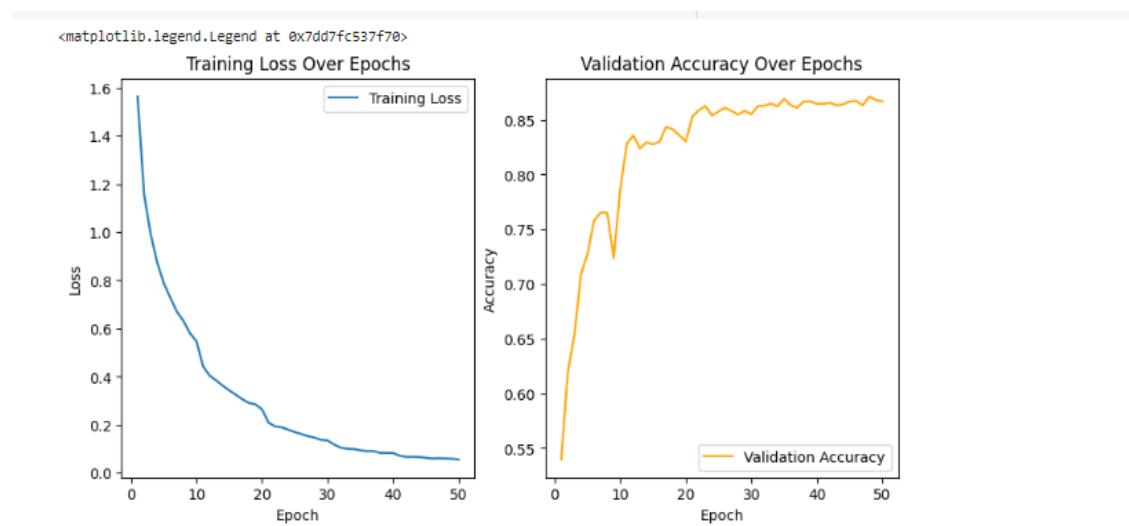
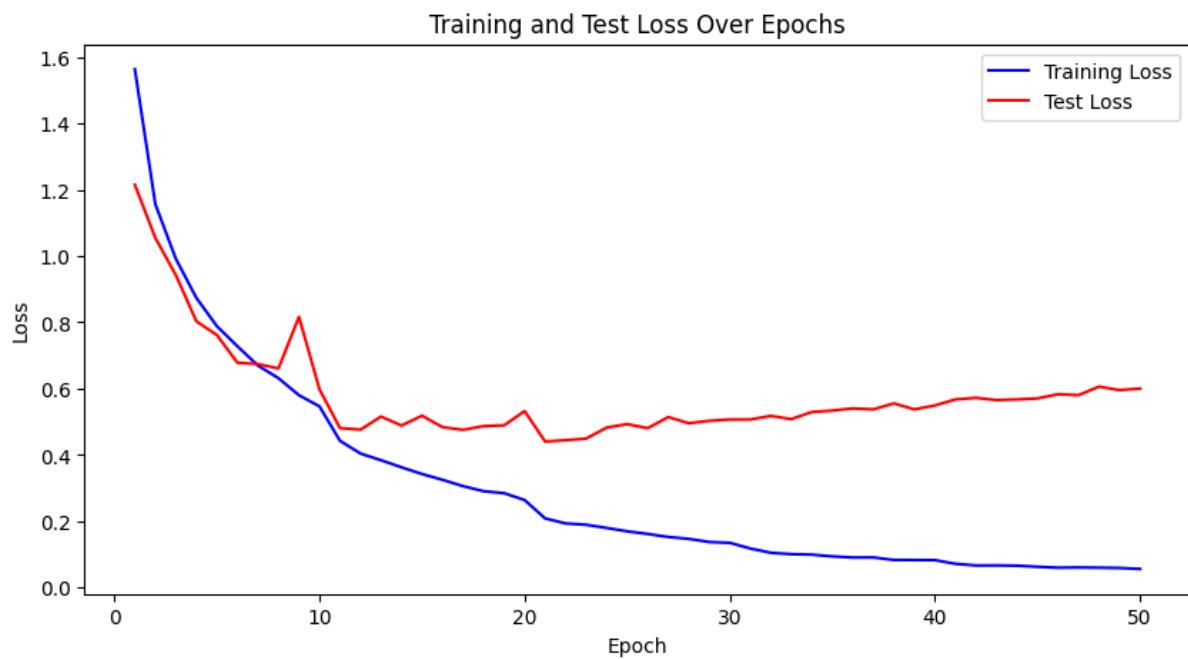
```
class ImprovedNet(nn.Module):  
    def __init__(self, input_channels=3, input_size=32):  
        super(ImprovedNet, self).__init__()  
        self.conv1 = nn.Conv2d(input_channels, 64, 3, padding=1)  
        self.bn1 = nn.BatchNorm2d(64)  
        self.conv2 = nn.Conv2d(64, 128, 3, padding=1)  
        self.bn2 = nn.BatchNorm2d(128)  
        self.conv3 = nn.Conv2d(128, 256, 3, padding=1)  
        self.bn3 = nn.BatchNorm2d(256)  
        self.conv4 = nn.Conv2d(256, 512, 3, padding=1)  
        self.bn4 = nn.BatchNorm2d(512)  
        self.pool = nn.MaxPool2d(2, 2)  
        fc_input_size = 512 * (input_size // 16) * (input_size // 16)  
  
        self.fc1 = nn.Linear(fc_input_size, 1024)  
        self.fc2 = nn.Linear(1024, 512)  
        self.fc3 = nn.Linear(512, 10)  
        self.dropout = nn.Dropout(0.5)  
  
    def forward(self, x):  
        x = self.pool(F.relu(self.bn1(self.conv1(x))))  
        x = self.pool(F.relu(self.bn2(self.conv2(x))))  
        x = self.pool(F.relu(self.bn3(self.conv3(x))))  
        x = self.pool(F.relu(self.bn4(self.conv4(x))))  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = self.dropout(x)  
        x = F.relu(self.fc2(x))  
        x = self.dropout(x)  
        x = self.fc3(x)  
        return x
```

## Confusion Matrix and Classification Report :

		Confusion Matrix									
		0	1	2	3	4	5	6	7	8	9
Actual	0	41974	554	1275	794	763	185	294	523	1975	1263
	1	485	47496	183	309	97	110	201	77	802	2440
	2	2730	145	37107	2097	2551	1722	1740	1007	286	165
	3	728	240	2190	33955	1822	7991	2738	1222	502	462
	4	645	72	1698	1899	38942	1086	1400	1965	152	91
	5	254	151	1532	6562	1632	39321	889	2013	169	277
	6	209	196	1401	1796	962	751	43909	180	194	152
	7	397	79	759	1451	1809	1504	196	42237	159	409
	8	1914	1091	254	439	117	115	194	132	42448	696
	9	794	2196	132	282	99	207	175	259	652	45104
		Predicted									

Classification Report:					
	precision	recall	f1-score	support	
0	0.84	0.85	0.84	49600	
1	0.91	0.91	0.91	52200	
2	0.80	0.75	0.77	49550	
3	0.68	0.65	0.67	51850	
4	0.80	0.81	0.81	47950	
5	0.74	0.74	0.74	52800	
6	0.85	0.88	0.87	49750	
7	0.85	0.86	0.86	49000	
8	0.90	0.90	0.90	47400	
9	0.88	0.90	0.89	49900	
accuracy			0.82	500000	
macro avg	0.82	0.83	0.83	500000	
weighted avg	0.82	0.82	0.82	500000	

## Evaluation Graphs:



### The Third Model :

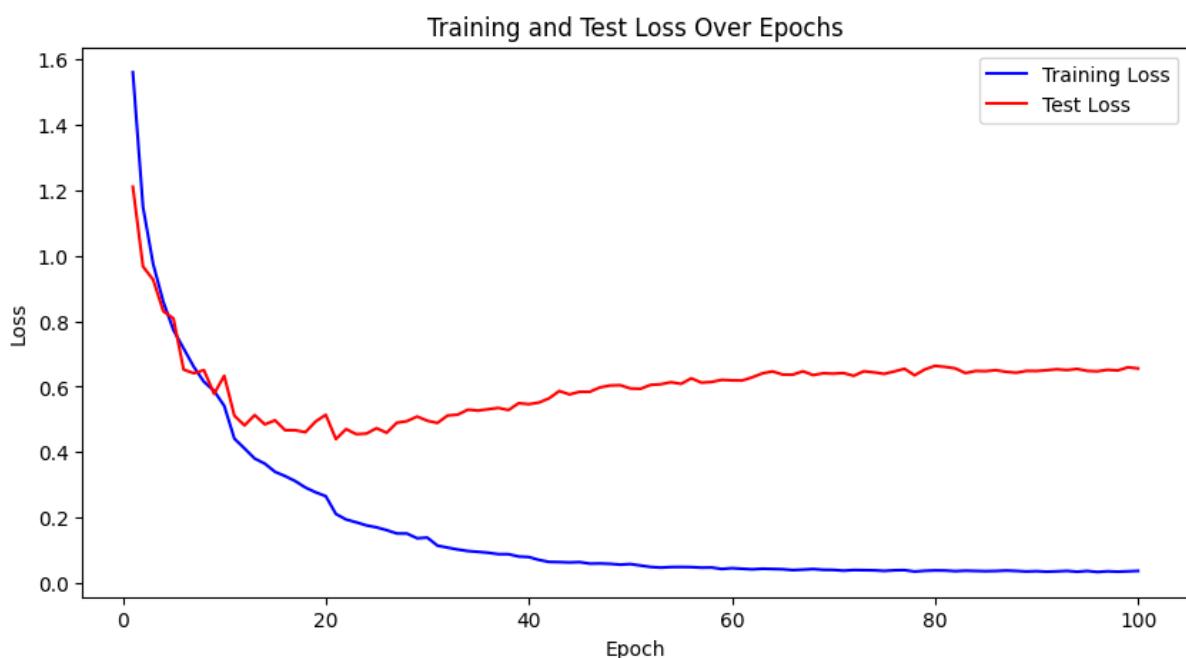
Same architecture but I increased the number of epochs from 50 to 100

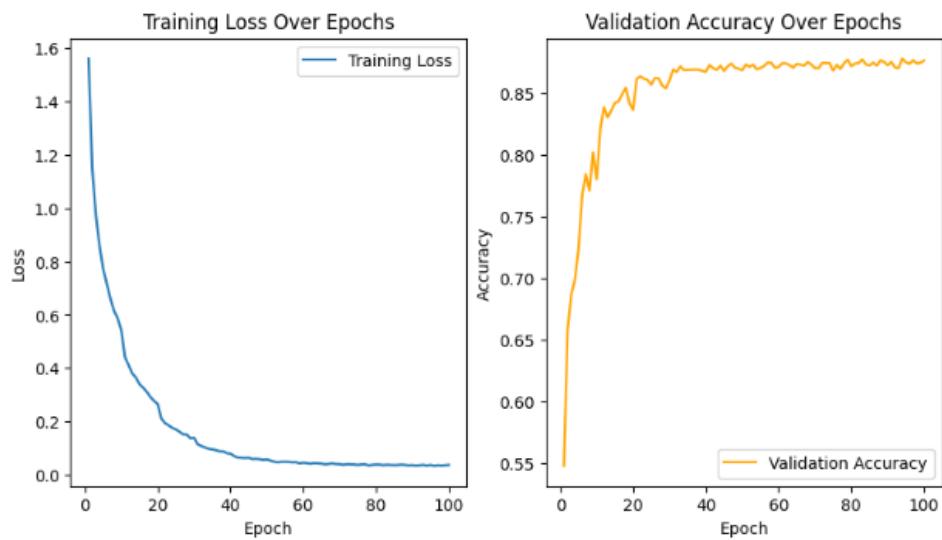
### Confusion Matrix and Classification Report :

		Confusion Matrix									
		0	1	2	3	4	5	6	7	8	9
Actual	0	91506	626	2573	1347	954	393	480	760	3206	1755
	1	861	89723	258	586	25	183	353	89	1195	3727
	2	4645	136	76565	3394	3837	2695	2843	1739	465	181
	3	1515	305	3747	69512	3220	14302	4265	1695	740	899
	4	1029	48	3793	3025	85871	2207	1487	3017	339	184
	5	416	282	2668	13019	2660	79663	1248	2907	264	473
	6	561	276	2433	3267	1555	1264	87388	263	241	252
	7	875	85	1608	1956	2860	3407	218	89027	188	676
	8	3090	1072	783	760	306	117	406	172	91975	1219
	9	1193	4008	198	769	203	127	259	266	1447	91330

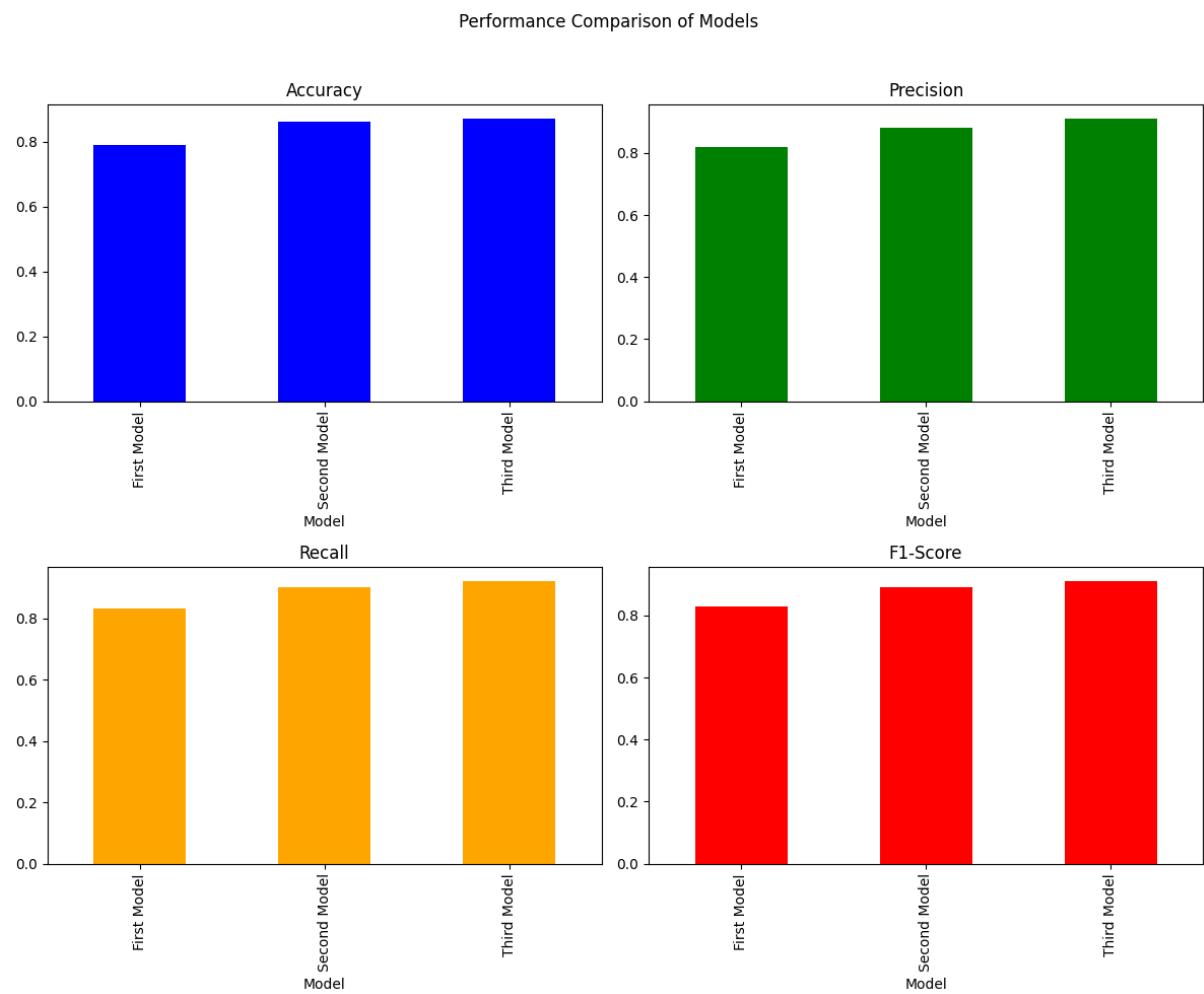
Classification Report:				
	precision	recall	f1-score	support
0	0.87	0.88	0.87	103600
1	0.93	0.92	0.93	97000
2	0.81	0.79	0.80	96500
3	0.71	0.69	0.70	100200
4	0.85	0.85	0.85	101000
5	0.76	0.77	0.77	103600
6	0.88	0.90	0.89	97500
7	0.89	0.88	0.89	100900
8	0.92	0.92	0.92	99900
9	0.91	0.92	0.91	99800
accuracy			0.85	1000000
macro avg	0.85	0.85	0.85	1000000
weighted avg	0.85	0.85	0.85	1000000

## Evaluation Graphs:





## Model comparison :



## Demo :

